

Nurse Notes — Judge Pitch

Aotearoa AI Hackathon Festival 2026 · University of Waikato

Read this first — three decisions baked into the script

1. This is a 5-minute pitch, not 3. The event's own judging criteria say "Clear, structured, within 5 minutes," and the official template is 0:30 Connect / 1:00 Problem / 2:00 Big Idea / 1:30 Impact . Your deck is *already* annotated to that template (slides 5 and 7 carry 2:00 and 1:30 in their eyebrow text). The 3-minute structure in the Phase 5 brief is generic hackathon advice — following it would cost you 40% of your allotted stage time and desync your slides. A 3-minute cutdown is at the end of this document if the room runs short on the day.

2. Your demo beat has changed — the old one no longer fires. CLAUDE.md §2 says you're deliberately keeping the 6/52 → "6 weeks and 5 days" bug as the moment the nurse catches an error. **That bug is fixed.** Commit 9c2a129 added src/lib/dates.js , which resolves every interval in JavaScript and hands the model finished dates. Commit 75d84f5 closed the "Activity limits" fabrication too. If you rehearse the old beat you'll be waiting on an error that never arrives.

The replacement is better, and it's *deterministic* — it fires every single time:

The model reaches "Activity limits", finds the source document silent on it, and writes **"Your notes do not mention this. [flag for nurse review]"** — then the nurse fills that gap on screen and approves.

A bug caught by luck says your model is unreliable. A model that reports what it doesn't know, into a field a human must clear, says **you designed the safety in**. That is a stronger claim and it survives contact with a live demo.

3. Never say "compliance" about a patient. NZ health research frames post-discharge failure as system-side — inadequate support, medication problems, unpreparedness. "Poor patient compliance" in front of NZ health judges will cost you the room. Say *"care a patient can't act on."* (Provider-side legal compliance under Right 5 is fine — that's about the hospital, not the patient.)

Speaking roles

Three speakers, five on stage. Every handover costs ~4 seconds and breaks rhythm — five speakers would burn 20 seconds of a 5-minute pitch on transitions alone.

Segment	Speaker	Note
Connect (0:30)	Himendra	Your story. See the gap below.
Problem (1:00)	Speaker B	Takes the baton straight from the story.
Big Idea + Demo (2:00)	Speaker C driving, B narrating	C's hands never leave the keyboard.
Impact + Ask (1:30)	Himendra (closes)	Same voice opens and closes — bookends the story.

Whoever isn't speaking: face the judges, don't watch the screen.

THE SCRIPT

0:00 – 0:30 · CONNECT — *Himendra's story*

Slide 1 (title)

YOUR STORY GOES HERE

Budget: 30 seconds ≈ 70–75 spoken words. Time it. This segment overrunning is the single most common way a hackathon pitch loses its demo.

What this slot has to do (per the official template: "*who you are, what you care about*"):

- Put one real person in the judges' heads — not a statistic.
- End on the moment of **not understanding**: the form signed, the instructions that made no sense, the question nobody asked.
- Do **not** mention the product. Not once. The product arrives at 1:30.

Land on this exact handover line so Speaker B can pick it up cleanly:

"...and they are not unusual. They are the majority."

(If your story ends somewhere else, tell B the last five words in advance so the baton pass is clean. Rehearse the seam, not just the halves.)

0:30 – 1:30 · PROBLEM — *Speaker B*

Slide 3 (the three statistics), then Slide 4 (Right 5)

Fifty-six percent of New Zealand adults — over 1.6 million people — can't obtain and understand basic health information well enough to make an informed decision. For Pacific people over fifteen, it's around ninety percent.

So we hand someone a discharge summary written clinician-to-clinician. Dense with shorthand — TDS, SOB, melaena, 6/52 — at the exact moment they're tired, unwell, and least able to decode it. Then sixteen percent of over-65s are back in hospital within thirty days. Higher for Māori, Pacific, and deprived areas.

[SLIDE 4] And here's the part that changed how we saw this. This isn't a service improvement. It's a legal right. Right 5 of the Code of Health and Disability Services Consumers' Rights guarantees every consumer effective communication — and I'll quote it — "*in a form, language, and manner that enables the consumer to understand.*"

Read that plainly. A consent form the patient didn't understand is arguably not compliant. Which means this isn't a nice-to-have that a hospital gets around to. It's an obligation that is currently, quietly, unmet — about a million and a half times over.

(≈150 words. Deliver Right 5 slowly — it's the sentence you want a judge to write down. Pause a full beat after "not compliant.")

1:30 – 3:30 · BIG IDEA + LIVE DEMO — C drives, B narrates

Slide 5, then switch to the app

[B] So we built Nurse Notes. It takes the clinical document, rewrites it in plain language, and puts a nurse between the model and the patient. Every time. No exceptions. And it runs entirely on this laptop — no cloud, no account, nothing uploaded.

[SWITCH TO APP — C clicks "Load sample"]

[B] This is a synthetic discharge summary — we've never used a real patient record, and we won't. On the left, the original. That badge says **61** — that's the count of pieces of clinical shorthand in this document that a patient cannot decode.

[C clicks "Generate" — text streams in]

[B] That's a four-billion-parameter model running locally. If we pulled the network cable right now, nothing would change.

Two things to watch as it writes.

First — the dates. This document has five slash intervals. 6/52 is six weeks. 5/7 is a five-day course. 4/7 is how long a cough lasted *before* admission. We don't ask the model to work those out — small models get it wrong, and we watched ours turn six weeks into "six weeks and five days." So we resolve them in JavaScript first, anchor them to the discharge date, and hand the model finished calendar dates to copy. It also has to tell an appointment apart from a course length — because putting a future date on a cough would be actively misleading. **[point]** There: a real date the patient can put in a calendar.

Second — this. [point at "Activity limits"] The notes say nothing about activity limits. So the model doesn't guess, and it doesn't reassure. It writes *"Your notes do not mention this — flag for nurse review."* A silent document is not a document that says no. Only a human can close that gap —

[C clicks into the field and types a correction]

— and that's the nurse doing it. **[C types name, clicks "Approve for release"]** Locked. Timestamped. Signed. **[C clicks "View on patient's phone"]** And this is what the patient walks out with — sixth-grade reading level, read-aloud, and a QR code that carries the summary *inside the code itself*. No app, no account, nothing fetched.

(≈300 words including pauses. Rehearse this on the presentation laptop until C never has to look for a button.)

3:30 – 5:00 · IMPACT + ASK — Himendra closes

Slide 7, then Slide 8

The study we built on took a discharge summary from eleventh-grade to sixth-grade reading level. Measured patient understanding went from thirteen percent to eighty-one percent. That's their number, not ours — and in that same study the model introduced errors. That's not a footnote we're burying. It's the reason the nurse gate exists, and it's why we showed you the gate instead of claiming the model is always right.

This maps to indicator 3.8.1 — coverage of essential health services. Care a patient can't act on isn't effective coverage. And to SDG 10.2, because low health literacy tracks with ethnicity and deprivation, so a document nobody can read isn't a usability problem — it's an equity problem.

On running local: that's a privacy property, and it's also a data-sovereignty one. Data that never leaves the ward can't be governed by someone else's jurisdiction. That matters for Māori data sovereignty, and it's why we didn't reach for a cloud API.

[SLIDE 8] What we need: a clinical partner to validate rewrites, verified te reo Māori and Pacific translations — human-verified, not machine-translated — and one pilot ward to measure whether comprehension actually moves.

Right 5 has been law for thirty years. The technology to actually meet it fits on a laptop and runs offline. **Our ask is simple: help us put a plain-language, clinician-approved form in front of the next patient who has to sign one.**

(≈240 words. Slow down on the last two sentences. Stop talking. Don't fill the silence.)

The three judge questions you will get

Q1. "How is this different from ChatGPT, or from Epic's Emmie?"

This is the question. Health chatbots are the most crowded category at SDG hackathons — the judges will be testing whether you know that.

Three ways. Epic's Emmie answers patient questions inside MyChart — it has no shipped feature that rewrites a full inpatient discharge summary, and New Zealand's public hospitals don't run on Epic, so none of it reaches a patient at Waikato. Second, we're not a chat box — the output is a fixed structure with a mandatory human approval step, and nothing reaches a patient unsigned. Third, paste a patient record into ChatGPT and you've just sent it offshore. Ours never leaves the machine.

Q2. "The model will get things wrong. How is this safe?"

Do not defend the model. Agree immediately, then show the architecture.

It will, and we assume it. That's why nothing is auto-sent — a nurse signs every release. But we didn't stop at the human. Where we could make an error impossible, we did: all date arithmetic is done in JavaScript, not by the model, because that's where a small model reliably fails. And the prompt is built so that a gap in the source produces a flag rather than a guess — the failure mode is "I don't know," not a confident invention. The human catches what's left. Roughly: deterministic code for what must be exact, the model for language, a nurse for judgement.

Q3. "Would a nurse actually use this? Isn't reviewing it more work?"

The adoption question. Have a number and a boundary.

Reviewing a one-page plain-language summary is faster than writing one from scratch, which is what a good nurse does verbally today, unpaid and unrecorded. And it produces an artifact the patient takes home. We're not claiming zero cost — we're claiming the cost lands in the right place, and it's why our ask is a pilot ward with measured comprehension rather than a rollout. If a nurse won't use it, we'd want to know that in week one.

Bonus — if a sponsor asks why you integrated no sponsor API

Deliberately. Our privacy and Māori data-sovereignty claim is that patient text never leaves the device — a cloud API call would break the core promise of the product to win a side prize. We'd rather be honest about the trade.

Deck changes to make before you present

Slide	Change	Why
5	Delete the footer "Concept mockup — the live build is a working web app (see next slide)"	It's no longer a mockup. Apologising for a real build is the worst thing on the deck.
6	Retitle step 4 to "Language & format — roadmap"	Te reo / Pacific / audio isn't built. Read-aloud <i>is</i> (browser SpeechSynthesis). Claiming the rest as shipped is the one thing that could sink you on a follow-up question.
6	Add a fifth step or a footer line: "Dates resolved in code, not by the model"	Your strongest technical differentiator currently appears nowhere in the deck.
7	Add "in the study we build on" <i>visually</i> , not just spoken	The 13% → 81% figure is borrowed. Attribute it on the slide or a judge will assume you measured it.
8	Fine as-is	Criteria match the event's actual four (Inspiration / Technology / Design / Presentation).

Demo reliability checklist

Run through this on the **presentation device**, not your own laptop.

The day before

- ☐ LM Studio installed on the presentation laptop, `google/gemma-3n-e4b` downloaded
- ☐ Temperature **0.25** (Settings, not Sampling)
- ☐ Developer tab → **Start Server** → confirm `http://localhost:1234/v1` responds

- ☐ `npm run build` → confirm `dist/index.html` opens from disk and works
- ☐ **Screen-record the full demo end to end.** Save it on the desktop, one click away
- ☐ Confirm the "Activity limits" flag actually fires on the sample — this is your beat now. If the model's phrasing drifts, that's fine; the *flag* is what you point at

Thirty minutes before

- ☐ Model **loaded in LM Studio and warmed** — run one generation. A cold first-token wait is 20 seconds of dead air
- ☐ App already open at the review screen, sample **not** yet loaded
- ☐ Notifications off. Every other tab and app closed
- ☐ Screen brightness up, laptop on mains power, sleep disabled
- ☐ Recording file open in a background window, ready to alt-tab

If it breaks Switch to the recording **without apologising**. Say:

"I'll show you the recording — same run, made this morning."

Then keep narrating exactly as scripted. Never say "it worked earlier."

Network: irrelevant, and say so. "If the venue WiFi dies, this demo doesn't notice" is a line worth using out loud.

3-minute cutdown (only if the room runs short)

Time	Segment	What changes
0:00–0:20	Story	Cut to ~45 words. One person, one moment of not understanding
0:20–0:50	Problem	56% / 1.6M, then straight to Right 5. Drop the readmission stat
0:50–2:20	Demo	Load → Generate → skip the date explanation → the flag → Approve → patient phone
2:20–3:00	Impact + ask	13%→81% (attributed), 3.8.1 + 10.2, on-device, the ask

Cut in that order. **Never cut the flag-and-approve moment** — it's the whole thesis.

One-sentence version, for a corridor conversation

Nurse Notes rewrites hospital discharge summaries into plain language on-device, and no patient sees a word of it until a nurse has corrected and signed it — because in New Zealand, communication a patient can understand isn't a courtesy, it's Right 5.